

# California Housing price

## Downloading the data

When we get the data, we get a file in the tgz form  
So we need to extract tgz file

```
1 import os
2 import tarfile
3 from six.moves import urllib
4
5 HOUSING_PATH = "datasets\\housing"
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131
132
133
134
135
136
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160
161
162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208
209
210
211
212
213
214
215
216
217
218
219
220
221
222
223
224
225
226
227
228
229
230
231
232
233
234
235
236
237
238
239
240
241
242
243
244
245
246
247
248
249
250
251
252
253
254
255
256
257
258
259
260
261
262
263
264
265
266
267
268
269
270
271
272
273
274
275
276
277
278
279
280
281
282
283
284
285
286
287
288
289
290
291
292
293
294
295
296
297
298
299
300
301
302
303
304
305
306
307
308
309
310
311
312
313
314
315
316
317
318
319
320
321
322
323
324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341
342
343
344
345
346
347
348
349
350
351
352
353
354
355
356
357
358
359
360
361
362
363
364
365
366
367
368
369
370
371
372
373
374
375
376
377
378
379
380
381
382
383
384
385
386
387
388
389
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407
408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431
432
433
434
435
436
437
438
439
440
441
442
443
444
445
446
447
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464
465
466
467
468
469
470
471
472
473
474
475
476
477
478
479
480
481
482
483
484
485
486
487
488
489
490
491
492
493
494
495
496
497
498
499
500
501
502
503
504
505
506
507
508
509
510
511
512
513
514
515
516
517
518
519
520
521
522
523
524
525
526
527
528
529
530
531
532
533
534
535
536
537
538
539
540
541
542
543
544
545
546
547
548
549
550
551
552
553
554
555
556
557
558
559
560
561
562
563
564
565
566
567
568
569
570
571
572
573
574
575
576
577
578
579
580
581
582
583
584
585
586
587
588
589
590
591
592
593
594
595
596
597
598
599
600
601
602
603
604
605
606
607
608
609
610
611
612
613
614
615
616
617
618
619
620
621
622
623
624
625
626
627
628
629
630
631
632
633
634
635
636
637
638
639
640
641
642
643
644
645
646
647
648
649
650
651
652
653
654
655
656
657
658
659
660
661
662
663
664
665
666
667
668
669
670
671
672
673
674
675
676
677
678
679
680
681
682
683
684
685
686
687
688
689
690
691
692
693
694
695
696
697
698
699
700
701
702
703
704
705
706
707
708
709
710
711
712
713
714
715
716
717
718
719
720
721
722
723
724
725
726
727
728
729
730
731
732
733
734
735
736
737
738
739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754
755
756
757
758
759
760
761
762
763
764
765
766
767
768
769
770
771
772
773
774
775
776
777
778
779
780
781
782
783
784
785
786
787
788
789
790
791
792
793
794
795
796
797
798
799
800
801
802
803
804
805
806
807
808
809
810
811
812
813
814
815
816
817
818
819
820
821
822
823
824
825
826
827
828
829
830
831
832
833
834
835
836
837
838
839
840
841
842
843
844
845
846
847
848
849
850
851
852
853
854
855
856
857
858
859
860
861
862
863
864
865
866
867
868
869
870
871
872
873
874
875
876
877
878
879
880
881
882
883
884
885
886
887
888
889
890
891
892
893
894
895
896
897
898
899
900
901
902
903
904
905
906
907
908
909
910
911
912
913
914
915
916
917
918
919
920
921
922
923
924
925
926
927
928
929
930
931
932
933
934
935
936
937
938
939
940
941
942
943
944
945
946
947
948
949
950
951
952
953
954
955
956
957
958
959
960
961
962
963
964
965
966
967
968
969
970
971
972
973
974
975
976
977
978
979
980
981
982
983
984
985
986
987
988
989
990
991
992
993
994
995
996
997
998
999
1000
```

## Loading csv file

use pandas to load csv file

```
1 import pandas as pd
2 def load_housing_data(housing_path=HOUSING_PATH):
3     csv_path = os.path.join(housing_path, "housing.csv")
4     return pd.read_csv(csv_path)
```

```
1 housing = load_housing_data()
2 housing.head()
```

```
1 .dataframe tbody tr th {
2     vertical-align: top;
3 }
4
5 .dataframe thead th {
6     text-align: right;
7 }
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income
0	-122.23	37.88	41.0	880.0	129.0	322.0	126.0	8.3252
1	-122.22	37.86	21.0	7099.0	1106.0	2401.0	1138.0	8.3014
2	-122.24	37.85	52.0	1467.0	190.0	496.0	177.0	7.2574
3	-122.25	37.85	52.0	1274.0	235.0	558.0	219.0	5.6431
4	-122.25	37.85	52.0	1627.0	280.0	565.0	259.0	3.8462

```
1 housing.info()
```

```
1 <class 'pandas.core.frame.DataFrame'>
2 RangeIndex: 20640 entries, 0 to 20639
3 Data columns (total 10 columns):
4  #   Column                Non-Null Count  Dtype
5  ---  ---
6  0   longitude              20640 non-null  float64
7  1   latitude               20640 non-null  float64
8  2   housing_median_age     20640 non-null  float64
9  3   total_rooms            20640 non-null  float64
10  4   total_bedrooms         20433 non-null  float64
```

```

11 | 5   population      20640 non-null float64
12 | 6   households      20640 non-null float64
13 | 7   median_income   20640 non-null float64
14 | 8   median_house_value 20640 non-null float64
15 | 9   ocean_proximity 20640 non-null object
16 | dtypes: float64(9), object(1)
17 | memory usage: 1.6+ MB

```

- All the attributes are numerical except `ocean_proximity`.  
And it is probably a catogries attribute. And we can use `value_counts` to find out

```

1 | housing["ocean_proximity"].value_counts()

```

```

1 | <1H OCEAN      9136
2 | INLAND        6551
3 | NEAR OCEAN    2658
4 | NEAR BAY      2290
5 | ISLAND         5
6 | Name: ocean_proximity, dtype: int64

```

- And we can also look at the numerical attributes

```

1 | housing.describe()

```

```

1 | .dataframe tbody tr th {
2 |     vertical-align: top;
3 | }
4 |
5 | .dataframe thead th {
6 |     text-align: right;
7 | }

```

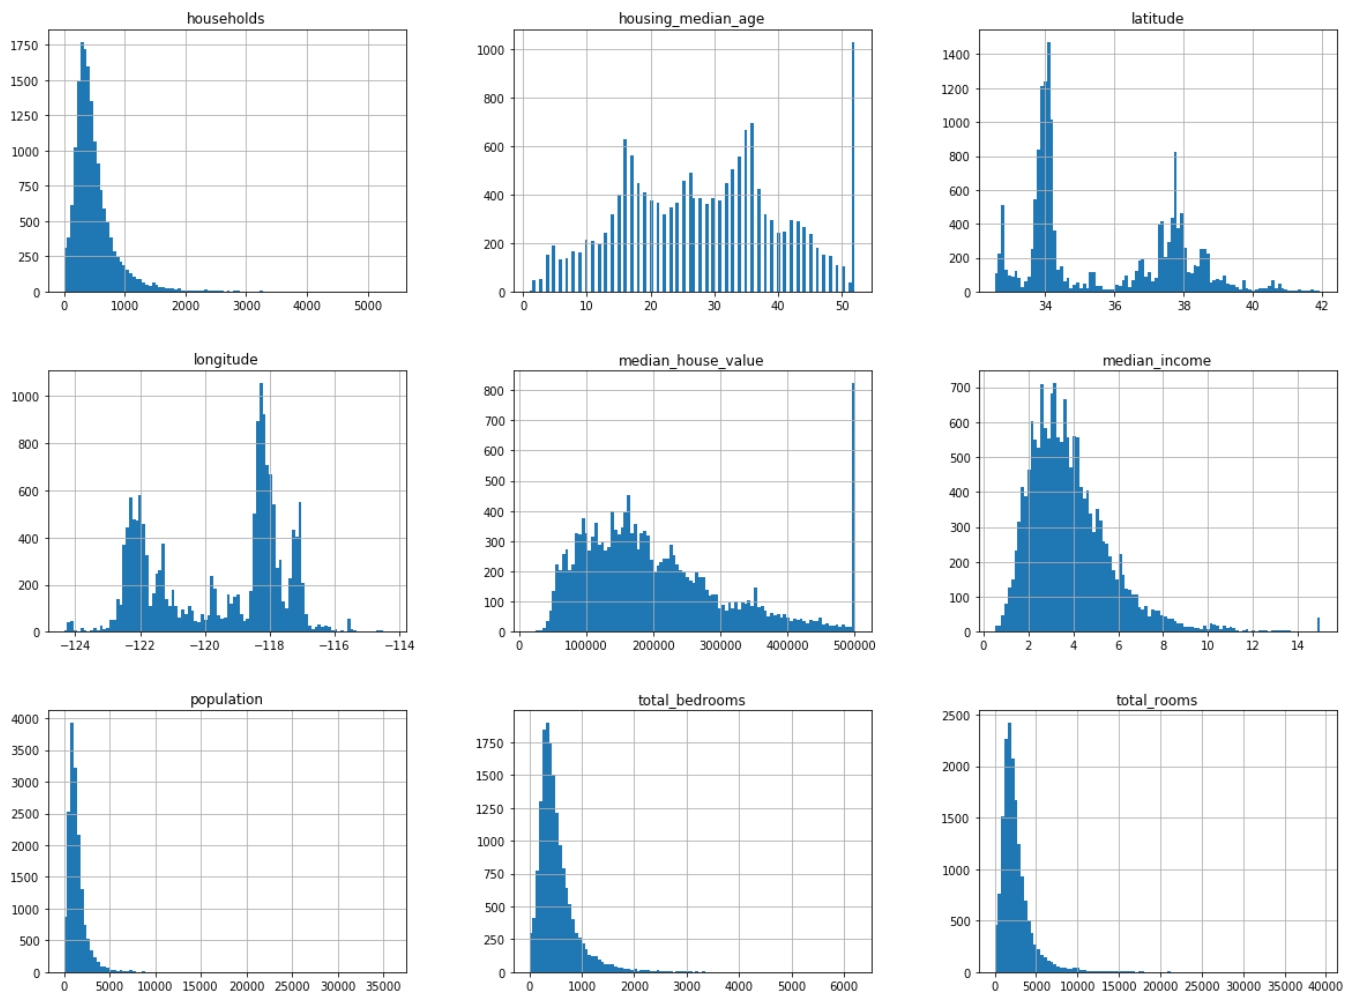
	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income
count	20640.000000	20640.000000	20640.000000	20640.000000	20433.000000	20640.000000	20640.000000	20640.000000
mean	-119.569704	35.631861	28.639486	2635.763081	537.870553	1425.476744	499.539680	3.870671
std	2.003532	2.135952	12.585558	2181.615252	421.385070	1132.462122	382.329753	1.899822
min	-124.350000	32.540000	1.000000	2.000000	1.000000	3.000000	1.000000	0.499900
25%	-121.800000	33.930000	18.000000	1447.750000	296.000000	787.000000	280.000000	2.563400
50%	-118.490000	34.260000	29.000000	2127.000000	435.000000	1166.000000	409.000000	3.534800
75%	-118.010000	37.710000	37.000000	3148.000000	647.000000	1725.000000	605.000000	4.743250
max	-114.310000	41.950000	52.000000	39320.000000	6445.000000	35682.000000	6082.000000	15.000100

we can also plot the histogram for each numerical attribute

```

1 | %matplotlib inline
2 | # only in jupyter notebook
3 | import matplotlib.pyplot as plt
4 | housing.hist(bins = 100, figsize=(20,15)) # figure size #bin : the number of the catogories
5 | plt.show()

```



## Create a Test Set

Create a test set: pick some examples randomly, approximately 20% of the dataset

```
1 import numpy as np
2
3 def split_train_test(data, test_ratio):
4     random_indices = np.random.permutation(len(data))
5     test_set_size = int( len(data) * test_ratio)
6     test_indices = random_indices[:test_set_size]
7     train_indices = random_indices[test_set_size:]
8     return data.iloc[train_indices], data.iloc[test_indices]
```

```
1 train_set, test_set = split_train_test(housing, 0.2)
2 print(len(train_set), "train +", len(test_set), "test", len(test_set)/len(housing))
```

```
1 16512 train + 4128 test 0.2
```

We got some problem using the method above  
that is when every time we apply the function, it generates different data

So we can generate random indices in the first time.

But if the dataset is updated?

One common way is using each instance's identifier to decide whether or not it should go in the test set

```
1 import hashlib
2
3 def test_set_check(identifier, test_ratio, hash):
4     return hash(np.int64(identifier)).digest()[-1] < 256 * test_ratio # the last byte of the md5 < 256 * ratio
5
6
7 def split_train_test_id(data, test_ratio, id_column, hash=hashlib.md5):
8     ids = data[id_column]
9     in_test_set = ids.apply(lambda id_: test_set_check(id_, test_ratio, hash)) # apply each id[i] to the function
10    # lambda : anonymous function
11    return data.loc[~in_test_set], data.loc[in_test_set]
```

Since the dataset do not have identifier column, we can set row index as ID:

```

1 housing_with_id = housing.reset_index() # add index column
2 housing_with_id["index"].head()

```

```

1 0    0
2 1    1
3 2    2
4 3    3
5 4    4
6 Name: index, dtype: int64

```

However, using index as identifier, we have to make sure that new data could only be added at the end of the dataset and no row ever get deleted.

Since the latitude and longitude for a city won't change for a short period, we can set them as identifier

```

1 housing_with_id["id"] = housing["longitude"] * 1000 + housing["latitude"]
2 housing_with_id["id"].head()
3 train_set, test_set = split_train_test_id(housing_with_id, 0.2, "id")

```

```

1 from sklearn.model_selection import train_test_split
2 train_set, test_set = train_test_split(housing, test_size=0.2, random_state=42) # done by skilt-learn

```

```


1 housing["median_income"].hist(bins = 15)

```

```

1 <matplotlib.axes._subplots.AxesSubplot at 0x1f15e6437f0>

```



To sample test set and training set according to the median income, we first generated categories through median\_income from each city

```

1 housing["income_cat"] = np.ceil(housing["median_income"] / 1.5)
2 housing["income_cat"].where(housing["income_cat"] < 5, 5.0, inplace=True)
3 # make categories which is greater than 5 merge into 5

```

Now using stratified sampling based on the income categories

```

1 from sklearn.model_selection import StratifiedShuffleSplit
2
3 split = StratifiedShuffleSplit(n_splits = 1, test_size = 0.2, random_state = 42)
4 for train_index, test_index in split.split(housing, housing["income_cat"]):
5     strat_train_set = housing.loc[train_index]
6     strat_test_set = housing.loc[test_index]
7
8 # To see the proportion of each categories
9 housing["income_cat"].value_counts() / len(housing)

```

```

1 3.0    0.350581
2 2.0    0.318847
3 4.0    0.176308
4 5.0    0.114438
5 1.0    0.039826
6 Name: income_cat, dtype: float64

```

Then remove column income\_categories

```

1 for set in (strat_train_set, strat_test_set):
2     set.drop(["income_cat"], axis = 1, inplace = True)

```

## Exporing the data and visualize the data


For not harming the training set, let's copy the data

```
1 housing = strat_train_set.copy()
```

### Visualizing geographical data


```
1 housing.plot(kind = "scatter", x = "longitude", y = "latitude") # scatter means discrete point
```

```
1 <matplotlib.axes._subplots.AxesSubplot at 0x1785b2ebbe0>
```



```
1 housing.plot(kind = "scatter", x = "longitude", y = "latitude", alpha=0.2,  
2           figsize = (8,6)) # alpha shows the density
```

```
1 <matplotlib.axes._subplots.AxesSubplot at 0x17856ba2be0>
```

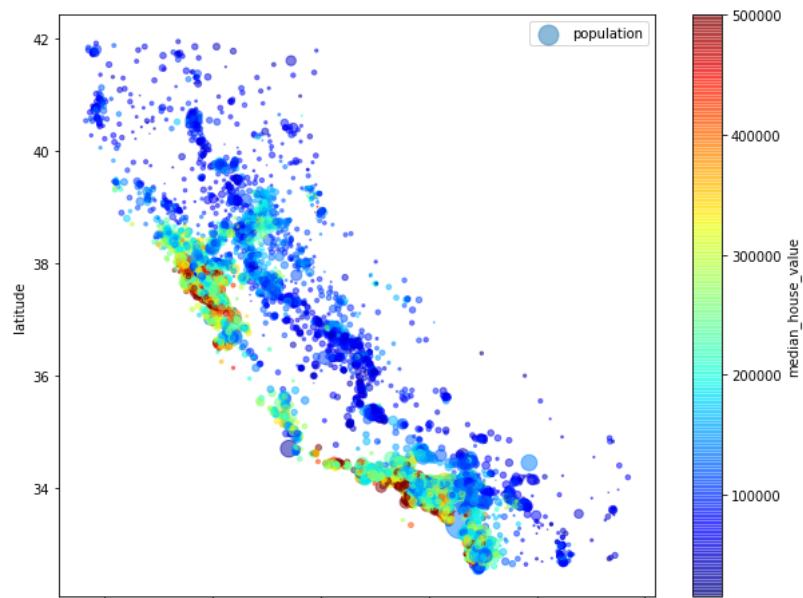


Now we are going to plot a figure that can show the density by the radius of each circle, price by the color

using predefined color map `jet` which ranges from blue to red

```
1 housing.plot(kind = "scatter", x = "longitude", y = "latitude", alpha = 0.5,  
2           s = housing["population"]/80, label = "population",  
3           c = "median_house_value", cmap = plt.get_cmap("jet"), colorbar = True,  
4           figsize = (10,8))  
5 plt.legend()
```

```
1 <matplotlib.legend.Legend at 0x1785baaf28>
```



## Looking for Correlations

```
1 | corr_matrix = housing.corr()
```

Looking at each attribute correlates with the median house value:

```
1 | corr_matrix["median_house_value"].sort_values(ascending = False)
```

```
1 | median_house_value    1.000000
2 | median_income         0.687160
3 | total_rooms           0.135097
4 | housing_median_age    0.114110
5 | households            0.064506
6 | total_bedrooms        0.047689
7 | population            -0.026920
8 | longitude             -0.047432
9 | latitude              -0.142724
10 | Name: median_house_value, dtype: float64
```

Another way to check correlation is check every numerical attribute against every other attribute.

There are 11 numerical attributes, we would get 121 plots which would not fit on a page. So we would use the data that most correlated with the median housing price

```
1 | housing.columns.tolist()
```

```
1 | ['longitude',
2 |  'latitude',
3 |  'housing_median_age',
4 |  'total_rooms',
5 |  'total_bedrooms',
6 |  'population',
7 |  'households',
8 |  'median_income',
9 |  'median_house_value',
10 |  'ocean_proximity']
```


```
1 | from pandas.plotting import scatter_matrix
2 |
3 | attributes = ["median_house_value", "median_income", "total_rooms", "housing_median_age"]
4 | scatter_matrix(housing[attributes], figsize=(12,8))
```

```
1 | array([[<matplotlib.axes._subplots.AxesSubplot object at 0x000001F15E3496D8>,
2 |        <matplotlib.axes._subplots.AxesSubplot object at 0x000001F15E453EF0>,
3 |        <matplotlib.axes._subplots.AxesSubplot object at 0x000001F15E42D390>,
4 |        <matplotlib.axes._subplots.AxesSubplot object at 0x000001F15E422908>],
5 |        [<matplotlib.axes._subplots.AxesSubplot object at 0x000001F15E66CEB8>,
```

```

6 <matplotlib.axes._subplots.AxesSubplot object at 0x000001F15E5C54A8>,
7 <matplotlib.axes._subplots.AxesSubplot object at 0x000001F15E999A58>,
8 <matplotlib.axes._subplots.AxesSubplot object at 0x000001F15ECC0080>],
9 [<matplotlib.axes._subplots.AxesSubplot object at 0x000001F15ECC00B8>,
10 <matplotlib.axes._subplots.AxesSubplot object at 0x000001F15E96ABA8>,
11 <matplotlib.axes._subplots.AxesSubplot object at 0x000001F15E555198>,
12 <matplotlib.axes._subplots.AxesSubplot object at 0x000001F15EC64630>],
13 [<matplotlib.axes._subplots.AxesSubplot object at 0x000001F15EC6CBA8>,
14 <matplotlib.axes._subplots.AxesSubplot object at 0x000001F15E12C198>,
15 <matplotlib.axes._subplots.AxesSubplot object at 0x000001F15E11B748>,
16 <matplotlib.axes._subplots.AxesSubplot object at 0x000001F15E1C7CF8>]],
17 dtype=object)


```



The most correlated attribute is median income, let's look closely

```
1 housing.plot(kind = "scatter", x = "median_income", y = "median_house_value", alpha = 0.2)
```

```
1 <matplotlib.axes._subplots.AxesSubplot at 0x1f15e1e1fd0>
```



## Attribute Combinations

the total number of rooms in a district is not very useful if we don't know how many households there are.

What we really want is the number of rooms per household.

Similarly, the total number of bedrooms by itself is not useful.

```

1 housing["rooms_per_household"] = housing["total_rooms"]/housing["households"]
2 housing["bedrooms_per_room"] = housing["total_bedrooms"]/housing["total_rooms"]
3 housing["population_per_household"] = housing["population"]/housing["households"]

```

Let's look the correlation of the new matrix

```

1 corr_matrix = housing.corr()
2 corr_matrix["median_house_value"].sort_values(ascending = False)

```

```

1 median_house_value      1.000000
2 median_income           0.687160
3 rooms_per_household     0.146285
4 total_rooms             0.135097
5 housing_median_age      0.114110
6 households              0.064506
7 total_bedrooms          0.047689
8 population_per_household -0.021985
9 population              -0.026920
10 longitude              -0.047432
11 latitude               -0.142724
12 bedrooms_per_room      -0.259984
13 Name: median_house_value, dtype: float64

```

## Prepare the Data for Machine Learning

First we copy the data, and let the data become labels and non-labels

```

1 housing = strat_train_set.drop("median_house_value", axis = 1)
2 housing_labels = strat_train_set["median_house_value"].copy()
3 housing.head()

```

```

1 .dataframe tbody tr th {
2     vertical-align: top;
3 }
4
5 .dataframe thead th {
6     text-align: right;
7 }

```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	ocean_proximity
<b>17606</b>	-121.89	37.29	38.0	1568.0	351.0	710.0	339.0	2.7042	<1H OCEAN
<b>18632</b>	-121.93	37.05	14.0	679.0	108.0	306.0	113.0	6.4214	<1H OCEAN
<b>14650</b>	-117.20	32.77	31.0	1952.0	471.0	936.0	462.0	2.8621	NEAR OCEAN
<b>3230</b>	-119.61	36.31	25.0	1847.0	371.0	1460.0	353.0	1.8839	INLAND
<b>3555</b>	-118.59	34.23	17.0	6592.0	1525.0	4459.0	1463.0	3.0347	<1H OCEAN

There are some missing values in total\_bedrooms, so they are three options:

- Get rid of the corresponding districts
- Get rid of the whole attribute
- Set the values to some value (zero, the mean, the median, etc.)

```

1 housing.dropna(subset=["total_bedrooms"]) # option 1
2 housing.drop("total_bedrooms", axis = 1) # option 2
3 median = housing["total_bedrooms"].median()
4 housing["total_bedrooms"].fillna(median) # option 3

```

```

1 17606 351.0
2 18632 108.0
3 14650 471.0
4 3230 371.0
5 3555 1525.0
6 ...
7 6563 236.0
8 12053 294.0
9 13908 872.0
10 11159 380.0
11 15775 682.0
12 Name: total_bedrooms, Length: 16512, dtype: float64

```

If option 3 is chosen, we need save `median`



sklearn apply a function Imputer instance

```
1 from sklearn.preprocessing import Imputer
2
3 imputer = Imputer(strategy = "median")
4
5 housing_num = housing.drop("ocean_proximity", axis = 1) # drop the non numerical attributes
```

```
1 C:\Users\Archibald Chain\AppData\Roaming\Python\Python37\site-packages\sklearn\utils\deprecation.py:66: DeprecationWarning: Class
  Imputer is deprecated; Imputer was deprecated in version 0.20 and will be removed in 0.22. Import impute.SimpleImputer from sklearn
  instead.
2 warnings.warn(msg, category=DeprecationWarning)
```

```
1 # fit the imputer instance to the training data using the fit() method:
2 imputer.fit(housing_num)
3 imputer.statistics_
4
```

```
1 array([-118.51 ,  34.26 ,  29.    , 2119.5   ,  433.    , 1164.    ,
2         408.    ,  3.5409])
```

```
1 X = imputer.transform(housing_num)
2 housing_tr = pd.DataFrame(X, columns = housing_num.columns)
```

## Handling Text and Categorical Attributes

Since the categories attributes are all in the 'string' type, we should turn every type into a categories.

```
1 from sklearn.preprocessing import LabelEncoder
2 encoder = LabelEncoder()
3 housing_cat = housing["ocean_proximity"]
4 housing_cat_encoded = encoder.fit_transform(housing_cat)
5 housing_cat_encoded
```

```
1 array([0, 0, 4, ..., 1, 0, 3])
```

```
1 encoder.classes_
```

```
1 array(['<1H OCEAN', 'INLAND', 'ISLAND', 'NEAR BAY', 'NEAR OCEAN'],
2        dtype=object)
```

However, the suitable solution should be giving 0 and 1 for each categories

```
1 from sklearn.preprocessing import OneHotEncoder
2 encoder = OneHotEncoder()
3 housing_cat_1hot = encoder.fit_transform(housing_cat_encoded.reshape(-1,1))
4 housing_cat_1hot
```

```
1 C:\Users\Archibald Chain\AppData\Roaming\Python\Python37\site-packages\sklearn\preprocessing\_encoders.py:415: FutureWarning: The
  handling of integer data will change in version 0.22. Currently, the categories are determined based on the range [0, max(values)],
  while in the future they will be determined based on the unique values.
2 If you want the future behaviour and silence this warning, you can specify "categories='auto'".
3 In case you used a LabelEncoder before this OneHotEncoder to convert the categories to integers, then you can now use the OneHotEncoder
  directly.
4 warnings.warn(msg, FutureWarning)
```

```
1 <16512x5 sparse matrix of type '<class 'numpy.float64'>'
2   with 16512 stored elements in Compressed Sparse Row format>
```

The output a Scipy sparse matrix, instead of Numpy array. This is useful when with thousands of categorical attributes with thousands of categories

```
1 housing_cat_1hot.toarray()
```

```
1 array([[1., 0., 0., 0., 0.],
2        [1., 0., 0., 0., 0.],
3        [0., 0., 0., 0., 1.],
4        ...,
5        [0., 1., 0., 0., 0.],
6        [1., 0., 0., 0., 0.],
7        [0., 0., 0., 1., 0.]])
```

- We can apply both transformations (from text categories to integer categories, then from integer categories to one-hot vectors) in one shot using the `LabelBinarizer` class:

```
1 from sklearn.preprocessing import LabelBinarizer
2 encoder = LabelBinarizer( sparse_output=True) # parameter true, it generate spares categories.
3 housing_cat_1hot = encoder.fit_transform(housing_cat)
4 housing_cat_1hot
```

```
1 <16512x5 sparse matrix of type '<class 'numpy.int32'>'
2   with 16512 stored elements in Compressed Sparse Row format>
```

## Custom Transformer

We can construct our own custom transformer

- with `TransformerMixin`, we get `fit_transformer()`
- with `BaseEstimator`, we get `set_parameter()` and `get_parameter()`

```
1 from sklearn.base import BaseEstimator, TransformerMixin
2
3 rooms_ix, bedrooms_ix, population_ix, household_ix = 3, 4, 5, 6
4 class CombinedAttributesAdder(BaseEstimator, TransformerMixin):
5     def __init__(self, add_bedrooms_per_room = True):
6         self.add_bedrooms_per_room = add_bedrooms_per_room
7     def fit(self, X, y = None):
8         return self
9     def transform(self, X, y = None): # we get a numpy matrix rather than a pandas DataFrame
10        rooms_per_household = X[:, 3]/ X[:, household_ix]
11        population_per_household = X[:, population_ix] / X[:, bedrooms_ix]
12        if self.add_bedrooms_per_room:
13            bedrooms_per_room = X[:, rooms_ix]/ X[:,bedrooms_ix]
14            return np.c_[X, rooms_per_household, population_per_household,
15                        bedrooms_per_room]
16        else:
17            return np.c_[X, rooms_per_household, population_per_household]
18
19 attr_adder = CombinedAttributesAdder(add_bedrooms_per_room=True)
20 housing_extra_attribs = attr_adder.transform(housing.values)
21 housing_extra_attribs
```

```
1 array([[ -121.89,  37.29,  38.0, ...,  4.625368731563422,  2.022792022792023,
2          4.467236467236467],
3        [ -121.93,  37.05,  14.0, ...,  6.008849557522124,  2.8333333333333335,
4          6.287037037037037],
5        [ -117.2,  32.77,  31.0, ...,  4.225108225108225,  1.9872611464968153,
6          4.144373673036093],
7        ...,
8        [ -116.4,  34.09,  9.0, ...,  6.34640522875817,  2.4059633027522938,
9          5.567660550458716],
10       [ -118.01,  33.82,  31.0, ...,  5.50561797752809,  3.568421052631579,
11          5.157894736842105],
12       [ -122.45,  37.77,  52.0, ...,  4.843505477308295,  1.8607038123167154,
13          4.538123167155425]], dtype=object)
```

```

1 class DataFrameSelector(BaseEstimator, TransformerMixin):
2     def __init__(self, attribute_names):
3         self.attribute_names = attribute_names
4     def fit(self, X, y = None):
5         return self
6     def transform(self, X):
7         return X[self.attribute_names].values

```

## Featuring Scaling

Two method to do it min-max scaling and standardization:

- Min-max scaling (many people call this normalization) is quite simple: values are shifted and rescaled so that they end up ranging from 0 to 1.
- Standardization is quite different: first it subtracts the mean value (so standardized values always have a zero mean), and then it divides by the variance so that the resulting distribution has unit variance.

## Transformation Pipeline

```

1 from sklearn.pipeline import Pipeline
2 from sklearn.preprocessing import StandardScaler
3 num_pipeline = Pipeline([
4     ('imputer', Imputer(strategy="median")),
5     ('attribs_adder', CombinedAttributesAdder()),
6     ('std_scaler', StandardScaler()),
7 ])
8 housing_num_tr = num_pipeline.fit_transform(housing_num)

```

```

1 C:\Users\Archibald Chain\AppData\Roaming\Python\Python37\site-packages\sklearn\utils\deprecation.py:66: DeprecationWarning: Class
Imputer is deprecated; Imputer was deprecated in version 0.20 and will be removed in 0.22. Import impute.SimpleImputer from sklearn
instead.
2     warnings.warn(msg, category=DeprecationWarning)

```

When call FeatureUnion's it runs each transformer's `transform()` method in parallel, waits for their output, and then concatenates them and returns the result

In the latest version of sklearn, there are some problems when using pipeline on LabelBinarizer

LabelBinarizer `fit_transform()` only have two parameters, while pipeline thought it has three

Solution given by stack over flow is add a custom LabelBinarizer

```

1 from sklearn.base import TransformerMixin #gives fit_transform method for free
2 class MyLabelBinarizer(TransformerMixin):
3     def __init__(self):
4         self.encoder = LabelBinarizer()
5     def fit(self, x, y=0):
6         return self
7     def transform(self, x, y=0):
8         return self.encoder.fit_transform(x)

```

However, using the LabelBinarizer above is still a problem.

We can using the following way to solve it

```

1 from sklearn.base import TransformerMixin #gives fit_transform method for free
2 class MyLabelBinarizer2(TransformerMixin):
3     def __init__(self, sparser_output = False):
4         self.encoder1 = LabelEncoder()
5         self.encoder2 = OneHotEncoder(categories='auto')
6         self.sparser_output = False
7     def fit(self, x, y=0):
8         return self
9     def transform(self, x, y=0):
10        encoded = self.encoder1.fit_transform(housing_cat)
11        if self.sparser_output:
12            return self.encoder2.fit_transform(encoded.reshape(-1,1))
13        else:
14            return self.encoder2.fit_transform(encoded.reshape(-1,1)).toarray()

```

```

1 from sklearn.pipeline import FeatureUnion
2
3 num_attribs = list(housing_num) # return string index of each column as a list
4 cat_attribs = ["ocean_proximity"]
5
6 num_pipeline = Pipeline([
7     ('selector', DataFrameSelector(num_attribs)),
8     ('imputer', Imputer(strategy="median")), #defined previous
9     ('attribs_adder', CombinedAttributesAdder()), # defined previous
10    ('std_scaler', StandardScaler())

```

```

11 ])
12 cat_pipeline = Pipeline([
13     ('selector', DataFrameSelector(num_attribs)),
14     ('label_binarizer', MyLabelBinarizer2())
15 ])
16 full_pipeline = FeatureUnion(transformer_list=[
17     ("num_pipeline", num_pipeline),
18     ("cat_pipeline", cat_pipeline)
19 ])

```

```

1 C:\Users\Archibald Chain\AppData\Roaming\Python\Python37\site-packages\sklearn\utils\deprecation.py:66: DeprecationWarning: Class
Imputer is deprecated; Imputer was deprecated in version 0.20 and will be removed in 0.22. Import impute.SimpleImputer from sklearn
instead.
2     warnings.warn(msg, category=DeprecationWarning)

```

```

1 # To apply the pipeline
2 housing_prepared = full_pipeline.fit_transform(housing)
3

```

```

1 C:\Users\Archibald Chain\AppData\Roaming\Python\Python37\site-packages\sklearn\preprocessing\_encoders.py:415: FutureWarning: The
handling of integer data will change in version 0.22. Currently, the categories are determined based on the range [0, max(values)],
while in the future they will be determined based on the unique values.
2 If you want the future behaviour and silence this warning, you can specify "categories='auto'".
3 In case you used a LabelEncoder before this OneHotEncoder to convert the categories to integers, then you can now use the OneHotEncoder
directly.
4     warnings.warn(msg, FutureWarning)

```

```

1 housing_prepared

```

```

1 array([[ -1.15604281,  0.77194962,  0.74333089, ...,  0.          ,
2          0.          ,  0.          ],
3        [ -1.17602483,  0.6596948 , -1.1653172 , ...,  0.          ,
4          0.          ,  0.          ],
5        [  1.18684903, -1.34218285,  0.18664186, ...,  0.          ,
6          0.          ,  1.          ],
7        ...,
8        [  1.58648943, -0.72478134, -1.56295222, ...,  0.          ,
9          0.          ,  0.          ],
10       [  0.78221312, -0.85106801,  0.18664186, ...,  0.          ,
11          0.          ,  0.          ],
12       [-1.43579109,  0.99645926,  1.85670895, ...,  0.          ,
13          1.          ,  0.          ]])

```

## Training and Evaluating on the Training Set

### Linear Regression Model

```

1 from sklearn.linear_model import LinearRegression
2
3 lin_reg = LinearRegression()
4 lin_reg.fit(housing_prepared, housing_labels)

```

```

1 LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)

```

Let's try it out on few instances

```

1 some_data = housing.iloc[:5]
2 some_data = np.array(list(some_data))
3 some_labels = housing_labels.iloc[:5]
4 some_data_prepared = full_pipeline.transform(some_data)
5 print("Predictions:\t", lin_reg.predict(some_data_prepared))
6 print("Labels:\t\t", list(some_labels))

```

```

1 C:\Users\Archibald Chain\AppData\Roaming\Python\Python37\site-packages\ipykernel_launcher.py:7: FutureWarning: Using a non-tuple
sequence for multidimensional indexing is deprecated; use 'arr[tuple(seq)]' instead of 'arr[seq]'. In the future this will be
interpreted as an array index, 'arr[np.array(seq)]', which will result either in an error or a different result.
2     import sys

```

```

1 -----
2
3 IndexError                                Traceback (most recent call last)
4
5 <ipython-input-219-343b2027d11d> in <module>
6     2 some_data = np.array(list(some_data))
7     3 some_labels = housing_labels.iloc[:5]
8 ----> 4 some_data_prepared = full_pipeline.transform(some_data)
9     5 print("Predictions:\t", lin_reg.predict(some_data_prepared))
10    6 print("Labels:\t\t", list(some_labels))

```

```

1 ~\AppData\Roaming\Python\Python37\site-packages\sklearn\pipeline.py in transform(self, X)
2     958         xs = Parallel(n_jobs=self.n_jobs)(
3         959             delayed(_transform_one)(trans, X, None, weight)
4 --> 960             for name, trans, weight in self._iter())
5     961         if not xs:
6     962             # All transformers are None

```

```

1 ~\AppData\Roaming\Python\Python37\site-packages\joblib\parallel.py in __call__(self, iterable)
2     919         # remaining jobs.
3     920         self._iterating = False
4 --> 921         if self.dispatch_one_batch(iterator):
5     922             self._iterating = self._original_iterator is not None
6     923

```

```

1 ~\AppData\Roaming\Python\Python37\site-packages\joblib\parallel.py in dispatch_one_batch(self, iterator)
2     757         return False
3     758     else:
4 --> 759         self._dispatch(tasks)
5     760         return True
6     761

```

```

1 ~\AppData\Roaming\Python\Python37\site-packages\joblib\parallel.py in _dispatch(self, batch)
2     714         with self._lock:
3     715             job_idx = len(self._jobs)
4 --> 716             job = self._backend.apply_async(batch, callback=cb)
5     717             # A job can complete so quickly than its callback is
6     718             # called before we get here, causing self._jobs to

```

```

1 ~\AppData\Roaming\Python\Python37\site-packages\joblib\_parallel_backends.py in apply_async(self, func, callback)
2     180     def apply_async(self, func, callback=None):
3     181         """Schedule a func to be run"""
4 --> 182         result = ImmediateResult(func)
5     183         if callback:
6     184             callback(result)

```

```

1 ~\AppData\Roaming\Python\Python37\site-packages\joblib\_parallel_backends.py in __init__(self, batch)
2     547         # Don't delay the application, to avoid keeping the input
3     548         # arguments in memory
4 --> 549         self.results = batch()
5     550
6     551     def get(self):

```

```

1 ~\AppData\Roaming\Python\Python37\site-packages\joblib\parallel.py in __call__(self)
2     223         with parallel_backend(self._backend, n_jobs=self._n_jobs):
3     224             return [func(*args, **kwargs)
4 --> 225                     for func, args, kwargs in self.items]
5     226
6     227     def __len__(self):

```

```

1 ~\AppData\Roaming\Python\Python37\site-packages\joblib\parallel.py in <listcomp>(.0)
2     223         with parallel_backend(self._backend, n_jobs=self._n_jobs):
3     224             return [func(*args, **kwargs)
4 --> 225                     for func, args, kwargs in self.items]
5     226
6     227     def __len__(self):

```

```

1 ~\AppData\Roaming\Python\Python37\site-packages\sklearn\pipeline.py in _transform_one(transformer, X, y, weight, **fit_params)
2     693
3     694 def _transform_one(transformer, X, y, weight, **fit_params):
4 --> 695     res = transformer.transform(X)
5     696     # if we have a weight for this transformer, multiply output
6     697     if weight is None:

```

```

1 ~\AppData\Roaming\Python\Python37\site-packages\sklearn\pipeline.py in _transform(self, X)
2     538         Xt = X
3     539         for _, _, transform in self._iter():
4 --> 540             Xt = transform.transform(Xt)
5     541         return Xt
6     542

```

```

1 <ipython-input-66-2a3b4f0f19d9> in transform(self, X)
2     5         return self
3     6     def transform(self, X):
4 ----> 7         return X[self.attribute_names].values

```

```

1 IndexError: only integers, slices (':',), ellipsis ('...'), numpy.newaxis ('None') and integer or boolean arrays are valid indices

```

```

1 from sklearn.metrics import mean_squared_error
2 housing_predictions = lin_reg.predict(housing_prepared)
3 lin_rmse = mean_squared_error(housing_labels, housing_predictions)
4 lin_rmse

```

```

1 4723628725.279794

```

## Decision Tree Regression

```

1 from sklearn.tree import DecisionTreeRegressor
2
3 tree_reg = DecisionTreeRegressor()
4 tree_reg.fit(housing_prepared, housing_labels)

```

```

1 DecisionTreeRegressor(criterion='mse', max_depth=None, max_features=None,
2                        max_leaf_nodes=None, min_impurity_decrease=0.0,
3                        min_impurity_split=None, min_samples_leaf=1,
4                        min_samples_split=2, min_weight_fraction_leaf=0.0,
5                        presort=False, random_state=None, splitter='best')

```

```

1 housing_predictions = tree_reg.predict(housing_prepared)
2 tree_mse = mean_squared_error(housing_labels, housing_predictions)
3 tree_rmse = np.sqrt(tree_mse)
4 tree_rmse

```

```

1 0.0

```

## Using Cross-Validation

We don't want to touch the test set until we are ready launch a model that we are very confident.

So we can use `train_test_split` to split the training set into a smaller training set and a validation set, then train the model against the smaller training set and evaluate them against the validation set.

K-fold cross-validation: it randomly splits the training set into 10 distinct subsets called folds, then train the Decision Tree Model 10 times. picking a different fold for evaluation every time and training on the other 9 folds. The result is an array containing 10 evaluation error:

```

1 from sklearn.model_selection import cross_val_score
2 tree_scores = cross_val_score(tree_reg, housing_prepared, housing_labels, scoring = "neg_mean_squared_error", cv = 10)
3 tree_rmse_scores = np.sqrt(-tree_scores)

```

```

1 def display_scores(scores):
2     print("Scores:", scores)
3     print("Mean:", scores.mean())
4     print("Standard deviation:", scores.std())
5
6 display_scores(tree_rmse_scores)

```

```

1 Scores: [69050.64876384 65660.63657415 75056.13197538 74405.72582172
2 75109.89210238 71037.39765743 69461.54545119 69014.78516735
3 74549.56823219 71807.09541249]
4 Mean: 71515.34271581142
5 Standard deviation: 3068.4773994700613

```

```

1 lin_scores = cross_val_score(lin_reg, housing_prepared, housing_labels, scoring = "neg_mean_squared_error", cv = 10)
2 lin_rmse_scores = np.sqrt(-lin_scores)
3 display_scores(lin_rmse_scores)

```

```

1 Scores: [67017.10226888 67116.66376981 68070.49102354 74774.53394933
2 68465.44788446 71478.52891582 65132.61122812 68494.45813594
3 71953.12890783 67969.53501985]
4 Mean: 69047.25011035803
5 Standard deviation: 2703.5105098669956

```

It seems the Linear Regression Model performs better than Decision Tree Model. That's true, The Decision Tree model is overfitting so badly.

## Random Forest Model

```

1 from sklearn.ensemble import RandomForestRegressor
2 forest_reg = RandomForestRegressor()
3 forest_reg.fit(housing_prepared, housing_labels)
4 forest_rmse = mean_squared_error(housing_labels, forest_reg.predict(housing_prepared))
5 forest_rmse = np.sqrt(forest_rmse)
6 forest_rmse

```

```

1 C:\Users\Archibald Chain\AppData\Roaming\Python\Python37\site-packages\sklearn\ensemble\forest.py:245: FutureWarning: The default value
of n_estimators will change from 10 in version 0.20 to 100 in 0.22.
2 "10 in version 0.20 to 100 in 0.22.", FutureWarning)

```

```

1 22826.619828397306

```

```

1 forest_scores = cross_val_score(forest_reg, housing_prepared, housing_labels, scoring = "neg_mean_squared_error", cv = 10)
2 forest_rmse_scores = np.sqrt(-forest_scores)

```

```

1 array([52886.91087803, 50029.92959489, 52265.17651867, 55219.16916582,
2 53784.50474138, 56045.76721275, 51640.28748997, 51755.10523981,
3 55944.04380411, 53207.74971784])

```

```

1 display_scores(forest_rmse_scores)

```

```

1 Scores: [52886.91087803 50029.92959489 52265.17651867 55219.16916582
2 53784.50474138 56045.76721275 51640.28748997 51755.10523981
3 55944.04380411 53207.74971784]
4 Mean: 53277.864436327225
5 Standard deviation: 1884.8867018860176

```

## Now we need to save the model

```

1 from sklearn.externals import joblib
2
3 joblib.dump(forest_reg, "my_model.pkl")

```

```

1 ['my_model.pkl']

```

```

1 forest_reg = joblib.load("my_model.pkl")

```

## Fine-Tune Model

## Grid Search

```
1 from sklearn.model_selection import GridSearchCV
2
3 param_grid = [
4     {'n_estimators':[3,10,30], 'max_features':[2,4,6, 8]},
5     {'bootstrap':[False], 'n_estimators':[3,10], 'max_features':[2,3,4]},
6 ]
7
8 forest_reg = RandomForestRegressor()
9 grid_search = GridSearchCV(forest_reg, param_grid, cv = 5, scoring = 'neg_mean_squared_error')
10
11 grid_search.fit(housing_prepared, housing_labels)
```

```
1 GridSearchCV(cv=5, error_score='raise-deprecating',
2             estimator=RandomForestRegressor(bootstrap=True, criterion='mse',
3                                             max_depth=None,
4                                             max_features='auto',
5                                             max_leaf_nodes=None,
6                                             min_impurity_decrease=0.0,
7                                             min_impurity_split=None,
8                                             min_samples_leaf=1,
9                                             min_samples_split=2,
10                                            min_weight_fraction_leaf=0.0,
11                                            n_estimators='warn', n_jobs=None,
12                                            oob_score=False, random_state=None,
13                                            verbose=0, warm_start=False),
14            iid='warn', n_jobs=None,
15            param_grid=[{'max_features': [2, 4, 6, 8],
16                           'n_estimators': [3, 10, 30]},
17                        {'bootstrap': [False], 'max_features': [2, 3, 4],
18                           'n_estimators': [3, 10]}],
19            pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
20            scoring='neg_mean_squared_error', verbose=0)
```

```
1 final_model = grid_search.best_estimator_
```

### Evaluate system on Test Set

```
1 x_test
```

```
1 .dataframe tbody tr th {
2     vertical-align: top;
3 }
4
5 .dataframe thead th {
6     text-align: right;
7 }
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	ocean_proximity
5241	-118.39	34.12	29.0	6447.0	1012.0	2184.0	960.0	8.2816	<1H OCEAN
10970	-117.86	33.77	39.0	4159.0	655.0	1669.0	651.0	4.6111	<1H OCEAN
20351	-119.05	34.21	27.0	4357.0	926.0	2110.0	876.0	3.0119	<1H OCEAN
6568	-118.15	34.20	52.0	1786.0	306.0	1018.0	322.0	4.1518	INLAND
13285	-117.68	34.07	32.0	1775.0	314.0	1067.0	302.0	4.0375	INLAND
...	...	...	...	...	...	...	...	...	...
20519	-121.53	38.58	33.0	4988.0	1169.0	2414.0	1075.0	1.9728	INLAND
17430	-120.44	34.65	30.0	2265.0	512.0	1402.0	471.0	1.9750	NEAR OCEAN
4019	-118.49	34.18	31.0	3073.0	674.0	1486.0	684.0	4.8984	<1H OCEAN
12107	-117.32	33.99	27.0	5464.0	850.0	2400.0	836.0	4.7110	INLAND
2398	-118.91	36.79	19.0	1616.0	324.0	187.0	80.0	3.7857	INLAND

4128 rows × 9 columns



```
1 | X_test = strat_test_set.drop("median_house_value", axis = 1)
2 | X_test_cat = X_test["ocean_proximity"]
3 | X_test_num = X_test.drop("ocean_proximity", axis = 1)
4 |
5 | X_test_num_array = num_pipeline.fit_transform(X_test_num)
6 |
7 | myLabelBinarizer = MyLabelBinarizer()
8 | X_test_cat_array = myLabelBinarizer.fit_transform(X_test_cat)
9 | y_test = strat_test_set["median_house_value"].copy()
10 |
11 |
12 | X_test_prepared = np.hstack((X_test_num_array, X_test_cat_array))
13 |
```

```
1 | (4128, 11)
```

```
1 | final_predictions = final_model.predict(X_test_prepared)
2 | final_mse = mean_squared_error(y_test, final_predictions)
3 | final_rmse = np.sqrt(final_mse)
```

```
1 | final_rmse
```

```
1 | 64692.9060362774
```